

an interesting decomposition

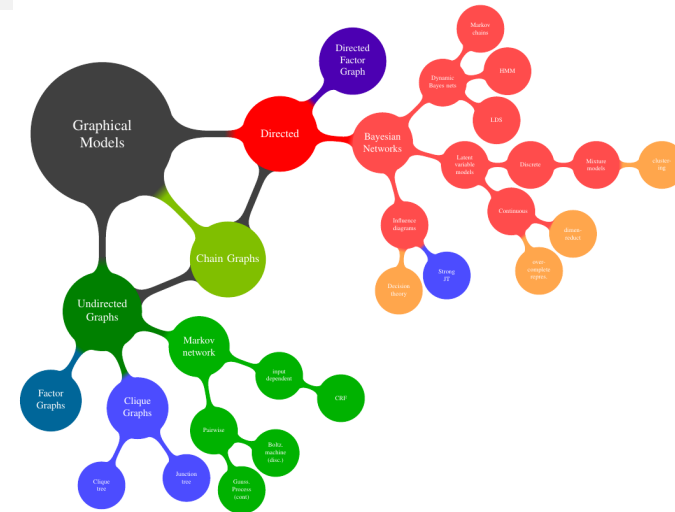
Big problems:

- 1 the learning problem:** given a training set  $\mathcal{D}$  of vectors  $\mathbf{x}$ , what is  $P(\mathbf{x} \mid \mathcal{D})$ ? ie. capture a complex joint distribution.
- 2 the inference problem:** given a training set  $\mathcal{D}$  of vectors  $\mathbf{x}$ , and knowledge of the values taken by some subset “obs” of the elements in  $x$  that are *observed*, what is  $P(x_i \mid \text{obs}, \mathcal{D})$ ? ie. find a posterior distribution.

answer: inference and learning are intractable in general, but perhaps could use prior knowledge of conditional independencies (eg. causation)?  $\Rightarrow$  PGMs


Today we will see that the LEARNING PROBLEM actually depends on a solution to the INFERENCE PROBLEM.

the PGM family



(from David Barber's amazing, free, book: google "Brml")  
(family-of-PGMs1.png in github week9)

the PGM family, by inference method



(from David Barber's amazing, free, book: google "Brml")  
(family-of-PGMs2.png in github week9)

## PGMs

I can think of three possible tasks:

- 1 infer  $p(x|\text{obs})$ , for any query variable  $x$  in the light of any set of observed variables  $\text{obs}$ .
- 2 infer the *most likely joint state*  $p(\mathbf{x}|\text{obs})$ , for all nodes simultaneously.
- 3 improve the tables (or learn from scratch) using a data set.

For a discrete-valued PGM...

#1 is solved by the **SUM-PRODUCT algorithm**, a.k.a. “probability propagation”, “belief propagation”, the “forward-backward algorithm”, and “turbo decoding”. (COMP307).

#2 is solved by the **MAX-SUM algorithm**, a.k.a. “Viterbi algorithm”. We will mention in the context of HMMs (next week).

#3 is the learning problem. Let's look at that...

## a very general view of learning: max likelihood

- my notation:  $\mathbf{x}$  for all variables  $(x_0, x_1, \dots)$ , but  $\mathbf{v}$  for those that are “visible” (observed),  $\mathbf{h}$  for those remaining hidden.  
 $\mathbf{x} = (\mathbf{v}, \mathbf{h})$ .
- we have a **model**, that has **parameters**  $\theta$
- For given  $\theta$  the model specifies a joint  $P(\mathbf{x} \mid \theta)$
- **data set** on the visibles:  $\mathcal{D} = \{\mathbf{v}_n\}, n = 1 \dots N$

For any given data set, the model gives a *likelihood*,  $P(\mathcal{D} \mid \theta)$ , often abbreviated as  $\mathcal{L}$ . This is the chance that our model could make the dataset *if we were to sample repeatedly from the joint*. It's a function of  $\theta$  so we can view it as a “surface” in  $\theta$ -space.

The most intuitive approach to learning: find and use the parameter values that **maximize the likelihood**. ie. those  $\theta$  for which our **data appears most likely**. This amounts to finding the highest point on a surface.

## unpacking the likelihood

$$\begin{aligned}\mathcal{L}(\theta) &= P(\mathcal{D} \mid \theta) \\ &= \prod_{\mathbf{v} \in \mathcal{D}} P(\mathbf{v} \mid \theta) \quad (\text{if data is i.i.d.})\end{aligned}$$

Nb. from now on we'll just drop the “ $\mid \theta$ ” from the r.h.s., as it occurs in everything.

$$\begin{aligned}\log \mathcal{L} &= \sum_{\mathbf{v} \in \mathcal{D}} \log P(\mathbf{v}) \\ &\propto \underbrace{\frac{1}{N} \sum_{\mathbf{v} \in \mathcal{D}} \log P(\mathbf{v})}_{\text{av. log likelihood per pattern}}\end{aligned}$$

## log likelihood of a dataset of $\mathbf{v}$

Consider a belief net, and a parameter  $\theta$  of (say) the  $j^{\text{th}}$  factor (which will be the factor  $P(x_j \mid \text{par}_j)$  in the net. We're going to look at the **gradient** of  $\log \mathcal{L}$  with respect to this parameter.

At the “best”  $\theta$ , this gradient must be zero, so this is a way of identifying the highest point.

**Here is a trick** for figuring out this gradient:

$$\text{1} \quad \frac{\partial}{\partial \theta} \log P = \frac{1}{P} \frac{\partial P}{\partial \theta} \quad \text{and thus (trivially)}$$

$$\text{2} \quad \frac{\partial P}{\partial \theta} = P \frac{\partial}{\partial \theta} \log P$$

## gradient of $\log \mathcal{L}$ for the whole data set

$$\log \mathcal{L} \propto \underbrace{\frac{1}{N} \sum_{\mathbf{v} \in \mathcal{D}} \log P(\mathbf{v})}_{\text{av. log likelihood per pattern}}$$

and so its gradient must be

$$\frac{\partial}{\partial \theta} \log \mathcal{L} \propto \frac{1}{N} \sum_{\mathbf{v} \in \mathcal{D}} \underbrace{\frac{\partial}{\partial \theta} \log P(\mathbf{v})}_{\text{so what's this?}}$$

## gradient of the log likelihood for a single pattern

$$\begin{aligned}
 \frac{\partial}{\partial \theta} \log P(\mathbf{v}) &= \frac{1}{P(\mathbf{v})} \frac{\partial}{\partial \theta} P(\mathbf{v}) && \leftarrow \text{via trick 1} \\
 &= \frac{1}{P(\mathbf{v})} \frac{\partial}{\partial \theta} \sum_{\mathbf{h}} P(\mathbf{v}, \mathbf{h}) && \leftarrow \text{sum rule} \\
 &= \sum_{\mathbf{h}} \frac{1}{P(\mathbf{v})} \frac{\partial}{\partial \theta} P(\mathbf{v}, \mathbf{h}) && \leftarrow \text{reordering} \\
 &= \sum_{\mathbf{h}} \frac{P(\mathbf{v}, \mathbf{h})}{P(\mathbf{v})} \frac{\partial}{\partial \theta} \log P(\mathbf{v}, \mathbf{h}) && \leftarrow \text{via trick 2} \\
 &= \underbrace{\sum_{\mathbf{h}} P(\mathbf{h} | \mathbf{v})}_{\text{av. over posterior!}} \frac{\partial}{\partial \theta} \log P(\mathbf{x}) && \leftarrow \text{product rule}
 \end{aligned}$$

## (back to) gradient of $\log \mathcal{L}$ for the whole data set

Now we can put the whole thing back together:

$$\frac{\partial}{\partial \theta} \log \mathcal{L} \propto \frac{1}{N} \underbrace{\sum_{\mathbf{v} \in \mathcal{D}}}_{\text{data}} \underbrace{\sum_{\mathbf{h}} P(\mathbf{h} | \mathbf{v})}_{\text{av. over posterior}} \frac{\partial}{\partial \theta} \log P(\mathbf{x})$$

Notice that now it's an average over the **the gradient of the joint**,  $\frac{\partial}{\partial \theta} \log P(\mathbf{x})$ , so that's a quantity of crucial importance!

EM algorithm:

- E: infer the posterior (holding  $\theta$  const.)
- M: take a step up the gradient (holding posterior const.)

eg: Mixtures of Gaussians, but EM applies to any directed PGM.





## directed PGMs = belief nets = "causal" nets

fully connected

naive Bayes

a chain

explaining away



$x_i \not\perp x_j$   
 $x_i \not\perp x_j \mid x_k$

$x_2 \not\perp x_3$   
 $x_2 \perp x_3 \mid x_1$

$x_1 \not\perp x_3$   
 $x_1 \perp x_3 \mid x_2$

$x_1 \perp x_2$ , and yet  
 $x_1 \not\perp x_2 \mid x_3$


## undirected PGMs (a.k.a. Markov random fields)

Graphical models describe joint probability distributions that *factor*. As we've seen, one way a distribution can factor is via application of the product rule to the joint as in, say,  $p(x, y, z) = p(x) p(y|x) p(z|x, y)$ , which corresponds to a directed graph called a Belief Net. However other factorisations exist. For example we could have

$$p(x, y, z) = \frac{1}{Z} \phi_A(x, y) \phi_B(y, z)$$

where  $Z$  is a normalisation factor. The  $\phi$  are usually called "potentials".

Eg. if  $x, y, z$  are binary,  $\phi_A$  and  $\phi_B$  are 2x2 tables.




The potentials  $\phi$  need only be positive.

One way to ensure this positivity is to use exponentials of another function:  $\phi_A = e^{E_A}$ . That way the function  $E$  is free to roam over *any* values. Then we have

$$p(x, y, z) = \frac{1}{Z} e^{E_A(x, y)} e^{E_B(y, z)} = \frac{1}{Z} e^{E_A(x, y) + E_B(y, z)}.$$

Physicists note: the  $E$  are completely analogous to (negative) energies in a physical system with Boltzmann distribution  $p$

Note that the potentials  $\phi$  *don't* need to be normalised along either their rows or columns.



Are  $x$  and  $z$  conditionally independent given  $y$ ?

$$p(x, z) = \sum_Y p(x, y, z) \propto \sum_Y \phi_A(x, y) \phi_B(y, z)$$

It seems clear that they won't de-couple if we don't know  $y$ . But:


$$\begin{aligned} p(x, z|y) &= \frac{p(x, y, z)}{\sum_x \sum_z p(x, y, z)} \propto \phi_A(x, y) \phi_B(y, z) \\ &= p(x|y) p(z|y) \end{aligned}$$

Once we know  $y$ , the distribution  $p(x, z|y)$  factors.

In an undirected graph, a variable becomes conditionally independent of *all other* variables, given its neighbours.


## PGM summary

### directed



- each factor is normalised
- product of all factors is automatically normalised
- can exhibit “explaining away”
- arrows are suggestive of a “causal” interpretation

### undirected



- factors aren't normalised
- product of all factors is not normalised
- no “causal” interpretation?
- seem to be a superset of directed models in fact...

Graphical models **simplify** the full joint by making **assumptions** about conditional independencies between variables.

## a general view of learning generative models

Earlier we looked at EM-like learning in directed graphical models. The steps were:

- 1 we wrote down the log likelihood,
- 2 took the gradient
- 3 discovered this involved adding up little gradients using samples from the data and the posterior over “hidden” nodes.
- 4 noted that we can use Gibbs Sampling to generate those samples

Now we will extend this to **any** graphical model, directed or not.

## a note on normalised vs unnormalised probabilities

Denote probabilities by  $P$ , or by  $P^*$  if they are not yet normalised.  
For example,

$$P(\mathbf{v}, \mathbf{h}) = \frac{P^*(\mathbf{v}, \mathbf{h})}{Z} \quad \text{with} \quad Z = \sum_{\mathbf{v}} \sum_{\mathbf{h}} P^*(\mathbf{v}, \mathbf{h}) \quad (1)$$

- In *some* cases of interest it is easy to ensure  $P$  is normalized (e.g. directed PGMs / belief nets).
- But in many cases it's easy to specify a plausible  $P^*$  yet hard to find  $Z$  and know  $P$  exactly (e.g. undirected PGMs)

the catch for undirected models

The sum in  $Z$  is over all *configurations* (all possible vectors  $\mathbf{x}$ ), so in general it's likely to be intractable.

## log likelihood of a dataset of $\mathbf{v}$

$$\begin{aligned} \log L &= \log P(\mathcal{D}) \\ &= \sum_{\mathbf{v} \in \mathcal{D}} \log P(\mathbf{v}) \\ &= \sum_{\mathbf{v} \in \mathcal{D}} \log (P^*(\mathbf{v})/Z) && \leftarrow \text{in terms of } P^* \\ &= \sum_{\mathbf{v} \in \mathcal{D}} (\log P^*(\mathbf{v}) - \log Z) \\ &\propto \underbrace{\frac{1}{N} \sum_{\mathbf{v} \in \mathcal{D}} \log P^*(\mathbf{v})}_{\text{av. log likelihood per pattern}} - \log Z \end{aligned}$$

Recap: **The trick** for finding the gradient of this: notice that

$$\begin{aligned} 1 \quad \frac{\partial}{\partial w} \log P &= (\frac{\partial}{\partial w} P)/P && \text{and conversely,} \\ 2 \quad \frac{\partial}{\partial w} P &= P \frac{\partial}{\partial w} \log P. \end{aligned}$$

## Recap: gradient of the first term (average of $\log P^*$ )

$$\begin{aligned} &\frac{\partial}{\partial w} \left[ \frac{1}{N} \sum_{\mathbf{v} \in \mathcal{D}} \log P^*(\mathbf{v}) \right] \\ &= \sum_{\mathbf{v} \in \mathcal{D}} \frac{1}{P^*(\mathbf{v})} \frac{\partial}{\partial w} P^*(\mathbf{v}) && \leftarrow \text{via trick 1} \\ &\vdots \\ &= \frac{1}{N} \sum_{\mathbf{v} \in \mathcal{D}} \underbrace{\sum_{\mathbf{h}} P^*(\mathbf{h} | \mathbf{v})}_{\text{av. over posterior}} \frac{\partial}{\partial w} \log P^*(\mathbf{x}) && \leftarrow \text{product rule} \end{aligned}$$

## gradient of the second term ( $\log Z$ )

The second term is all about the normalisation factor,  $Z$ .  
(NB. gradient will be *automatically* be zero in any belief net!)

$$\begin{aligned} \frac{\partial}{\partial w} \log Z &= \frac{1}{Z} \frac{\partial}{\partial w} \sum_{\mathbf{v}} \sum_{\mathbf{h}} P^*(\mathbf{v}, \mathbf{h}) && \leftarrow \text{trick 1} \\ &= \frac{1}{Z} \sum_{\mathbf{v}} \sum_{\mathbf{h}} \frac{\partial}{\partial w} P^*(\mathbf{v}, \mathbf{h}) \\ &= \frac{1}{Z} \sum_{\mathbf{v}} \sum_{\mathbf{h}} P^*(\mathbf{v}, \mathbf{h}) \frac{\partial}{\partial w} \log P^*(\mathbf{v}, \mathbf{h}) && \leftarrow \text{trick 2} \\ &= \underbrace{\sum_{\mathbf{v}} \sum_{\mathbf{h}} P(\mathbf{v}, \mathbf{h})}_{\text{average over joint!}} \frac{\partial}{\partial w} \log P^*(\mathbf{v}, \mathbf{h}) && \leftarrow \text{via eqtn 1} \end{aligned}$$

## gradient as a whole

Putting the two terms back together we get a total gradient of:

$$\frac{\partial}{\partial w} \log L \propto \underbrace{\frac{1}{N} \sum_{\mathbf{v} \in \mathcal{D}} \sum_{\mathbf{h}} P(\mathbf{h} | \mathbf{v})}_{\text{data} \quad \text{av. over posterior}} \frac{\partial}{\partial w} \log P^*(\mathbf{x}) - \underbrace{\sum_{\mathbf{v}, \mathbf{h}} P(\mathbf{v}, \mathbf{h})}_{\text{av. over joint}} \frac{\partial}{\partial w} \log P^*(\mathbf{x})$$

Both terms are some sort of average over  $\frac{\partial}{\partial w} \log P^*(\mathbf{x})$ , so that's a quantity of crucial importance.

Another way to write the overall gradient:

$$\left\langle \frac{\partial}{\partial w} \log P^*(\mathbf{x}) \right\rangle_{\mathbf{v} \in \mathcal{D}, \mathbf{h} \sim P(\mathbf{h} | \mathbf{v})} - \left\langle \frac{\partial}{\partial w} \log P^*(\mathbf{x}) \right\rangle_{\mathbf{x} \sim P(\mathbf{x})}$$

clamped / wake phase

↑↑↑ conditioned hypotheses

unclamped / sleep / free phase

↓↓↓ random fantasies